

Computer Architecture and Technology Area

Universidad Carlos III de Madrid





# OPERATING SYSTEMS

Lab 1. System calls

**BACHELOR'S DEGREE IN COMPUTER SCIENCE AND  
ENGINEERING**

Year 2018/2019

	<p align="center"> <b>Informatics Department</b>  <b>BACHELOR'S DEGREE IN COMPUTER</b>  <b>SCIENCE AND ENGINEERING</b>  <b>Operating Systems(2018-2019)</b>  <b>Lab 1 – System calls</b> </p>	
---	---	---

## Index

1. Lab Statement	3
1.1. Lab description	3
mycat	3
mys	4
mysize	5
1.2. Support code	7
1.3. Test suite	7
2. Assignment submission	8
2.1. Deadline	8
2.2. Submission	8
2.3. Files to be submitted	8
3. Rules	10
Appendix – System calls	11
I/O system calls	11
File related system calls	12
Manual (man command).	12
Bibliography	13

## 1. Lab Statement

This lab allows the student to familiarize with Operating System calls (especially related to the file system management) following the POSIX standard. Unix allows you to make calls directly to the Operating System from a program implemented in a high-level language, in particular, C language.

Most of input/output (I/O) operations in Unix can be done using uniquely five calls: `open`, `read`, `write`, `lseek` and `close`.

For the Operating System kernel, all files opened are identified using *file descriptors*. A file descriptor is a non negative integer. When we open a file that already exists, the kernel returns a file descriptor to the process. When we want to read or write from/to a file, we identify the file with the file descriptor that was returned by the `open` call.

Each open file has a *current read/write position* ("**current file offset**"). It is represented by a non-negative integer that measures the number of bytes from the beginning of the file. The read and write operations normally start at the current position and create an increment in that position equal to the number of bytes read or written. By default, this position is initialized to 0 when a file is opened, unless the option `O_APPEND` is specified. The current position (`current_offset`) of an open file can be changed explicitly using the system call `lseek`.

To manipulate directories, you can use the system calls *opendir*, *readdir* and *closedir*. An open directory is identified with a directory descriptor, which is a pointer of type `DIR` (*DIR\**). When we open a directory with *opendir*, the kernel returns a directory descriptor from which the different entries to that directory can be read using the calls to the function *readdir*. The call *readdir* returns a directory entry in a pointer to a structure of type `dirent` (*struct dirent\**). Such structure will contain the fields corresponding to that entry such as the name of the entry, or the type (if it is a normal file, another directory, symbolic links, etc.). Repeated calls to the function *readdir* will be returning the next entries in an open directory.

### 1.1. Lab description

In this lab you will be implementing three C programs which use the system calls that were previously described. These programs will be *mycat*, *myls* and *mysize*. For this purpose, you will have the following files with initial code *mycat.c*, *myls.c* and *mysize.c*.

## mycat

The first program, **mycat** shall open the file specified by argument and will show its content through standard output (the console) using the calls *open*, *read*, *write* and *close*. For this purpose:

- It will open (*open*) the file specified as parameter.
- It will read (*read*) the contents of the file using a 1024-byte intermediate buffer.
- It will write (*write*) the content of the buffer in the standard output. Use the constant *STDOUT\_FILENO* as value of the descriptor to write to the standard output.
- Finally, it will close (*close*) the descriptor.

Usage example:

```
./mycat p1_pruebas/f1.txt
```

Nombre1	V	32	09834320	24500
Nombre2	M	35	32478973	27456
Nombre3	V	53	98435834	45000

Usage: **./mycat** <path\_input\_file>

### Requirements :

- The program must show the whole contents of the file.
- The output must match exactly the output of an equivalent *cat* execution.
- The program must return -1 if no argument was passed.
- The program must return -1 if there was an error when opening the file (e.g. the file does not exist).

**Test suggestion<sup>1</sup>:** Check that the output of the program for one file corresponds with the one offered by the command *cat* (no extra arguments) over that same file.

## myls



The second program **myls**, will open a directory passed as parameter (or the current directory if no directory is specified) and print on the screen all the entries that this directory contains, one per line.

This program will:

- Obtain the specified directory from the arguments of the program or obtain the current directory using the call *getcwd*. Use the constant *PATH\_MAX* as maximum size that can have the path of the current directory.
- Open the directory using *opendir*.

---

<sup>1</sup> Passing this test is no guarantee of having the maximum mark in this exercise. It is just a suggestion so the students can check the general execution of their program. The students must also meet the other requirements of the program, write an adequate code, comment it, test extreme cases, and, in general, meet the other demands described in this statement.

	<p style="text-align: center;">Informatics Department BACHELOR'S DEGREE IN COMPUTER SCIENCE AND ENGINEERING Operating Systems(2018-2019)</p> <p style="text-align: center;">Lab 1 - System calls</p>	
---	--	---

- Then, it will read in each of the entries of the directory using *readdir* and print the name of the entry using *printf*.
- Finally, it will close the descriptor of the directory through the call *closedir*.

Usage example:

```
./mys p1_pruebas/
dirC
f1.txt
dirA
f2.txt
.
..
```

Usage: **./mys** <directory>

Usage 2: **./mys**

#### Requirements:

- The program must list all the entries in the directory, following the order in which the call to *readdir* returns them, and showing each entry in one line.
- The output must match exactly the output of an equivalent *ls* execution
- The program must list the entries of the directory passed as parameter (usage 1), or from the current directory if no parameter was passed.
- The program must return -1 if an error happened while opening the directory (e.g., the directory does not exist).

**Test suggestion<sup>2</sup>:** Check that the output of the program over a directory corresponds with the one obtained with the command **ls -f -l** over that same directory.

#### mysize

The third program, *mysize*, will obtain the current directory, and list all regular files contained in it, as well as their respective size in bytes. For this purpose:

- It will obtain the current directory using the call *getcwd*. Use the constant `PATH_MAX` as maximum size of the path of the current directory.
- Open the file using *opendir*.

---

<sup>2</sup> Passing this test is no guarantee of having the maximum mark in this exercise. It is just a suggestion so the students can check the general execution of their program. The students must also meet the other requirements of the program, write an adequate code, comment it, test extreme cases, and, in general, meet the other demands described in this statement.

- Then, it will read the entries of the directory using *readdir*.
- If the entry is a regular file (field *d\_type* from the structure *dirent* equal to the constant *DT\_REG*).
  - Open the file using *open*.
  - Move the file pointer to the end of the file and obtain its value with *lseek*.
  - Close the file with *close*.
  - Print the name of the file (field *d\_name* of the structure *dirent*), followed by a tab character, and the size obtained by *lseek*, ending with an EOF (end of line) character.
- This procedure will be repeated for every entry in the directory.
- Finally, it will close the directory descriptor with *closedir*.

Usage example:



```
cd p1_pruebas/

../mysize
f1.txt  87
f2.txt  87
```

**Usage: ./mysize**

#### **Requirements :**

- The program must show the name and size of all the regular files in the directory, following the order in which *readdir* returns them, and showing the data of each file in one line.
- The program will only show the data from regular files.
- The program will show the data using the following format: *<name><tab><size>*
- The program must return -1 if there was an error when opening the file.

	<p>Informatics Department BACHELOR'S DEGREE IN COMPUTER SCIENCE AND ENGINEERING Operating Systems(2018-2019)</p> <p>Lab 1 - System calls</p>	
---	--	---

## 1.2. Support code

Support code is provided in the file `p1_llamadas_2019.zip`. To extract its contents, you can use the **unzip** command:

```
unzip p1_llamadas_2019.zip
```

As a result, you will find a new directory `p1_llamadas/`, in which you must code the different programs. Inside this directory you will find:

### Makefile

File used by the make tool to compile all programs. **Do not modify this file.** Use `$ make` to compile the programs and `$ make clean` to remove the compiled files.

### mycat.c

C Source file to code *mycat*. **Must be modified by the student.**

### myls.c

C Source file to code *myls*. **Must be modified by the student.**

### mysize.c

Source file to code *mysize*. **Must be modified by the student.**

### p1\_tests/

This directory contains example files and directories to be able to execute and test your programs.

### Corrector\_ssoo\_p1\_2019.py

Basic test suite. Validates submission formal requirements and verifies basic functionality.

## 1.3. Test suite

The python script **corrector\_ssoo\_p1\_2019.py** should be used to verify that the student submission follows the format conventions (it has the correct names and it is well compressed). The tester must be executed in the Linux computers from the labs of the Computer Science and Engineering Department lab or in the **guernika.lab.inf.uc3m.es** server.

The command to execute the tester is the following:

```
python corrector_ssoo_p1_2019.py <submission_file.zip>
```

where *submission\_file.zip* is the file that it is going to be submitted to Aula Global (see next section). Example:

```
$ python corrector_ssoo_p1_2019.py ssoo_p1_100555555_100666666.zip
```

The tester will print on the console messages stating whether the required format is correct or not, and the results of basic functionality tests.

## 2. Assignment submission

### 2.1. *Deadline*

You must submit your work using Aula Global before the **3<sup>rd</sup> of March of 2019 at 23:55.**

### 2.2. *Submission*

The submission must be done using Aula Global using the links available in the first assignment section. The submission must be done separately for the code and report. The report will be submitted through the **TURNITIN** tool.

### 2.3. *Files to be submitted*



You must submit the code in a zip compressed file with name `ssoo_p1_AAAAAAAAAA_BBBBBBBBBB_CCCCCCCCC.zip` where A...A, B...B and C...C are the student identification numbers of the group. A maximum of three students is allowed per group. The file to be submitted must contain:

- o **Makefile**
- o **mycat.c**
- o **mysize.c**
- o **myls.c**

The report must be submitted as a PDF file. Notice that only PDF files will be evaluated. The file must be named `ssoo_p1_AAAAAAAAAA_BBBBBBBBBB_CCCCCCCCC.pdf`. The report must contain at least:

- **Title page** staging the name of the authors and their student identification numbers.
- **Table of contents.**
- **Description of the code** detailing the main functions it is composed of. Do not include any source code in the report, since it will be ignored.
- **Tests plan** detailing the objective of each test, the expected output, procedure and, optionally, obtained result with your implementation. Consider the following:
  - o Avoid duplicated tests that target the same code paths with equivalent input parameters.
  - o This section will be evaluated according to the coverage of the test plan, nor the overall number of tests.
  - o You can tabulate the information for clarity and concision.





	<p>Informatics Department BACHELOR'S DEGREE IN COMPUTER SCIENCE AND ENGINEERING Operating Systems(2018-2019) Lab 1 – System calls</p>	
---	---	---

- o Compiling without warnings does not guarantee that the program fulfills the requirements. Analyze your code to detect potential sources of errors and build your test plan accordingly.
- **Conclusions** describing the main problems found and how they have been solved. Additionally, you can include any personal conclusions derived from this assignment.

The PDF file must be submitted using the TURNITIN link. Do not neglect the quality of the report as it is a significant part of the grade of each assignment. Make sure that every page except the title page is numbered and the test is justified. **The report must not be longer than 8 pages including all the formerly detailed elements.**

**NOTE:** Only the last version of the submitted files will be reviewed.

	<p>Informatics Department BACHELOR'S DEGREE IN COMPUTER SCIENCE AND ENGINEERING Operating Systems(2018-2019) Lab 1 - System calls</p>	
---	---	---

### 3. Rules

- 1) **Programs that do not compile or do not meet the requirements will receive a mark of zero.**
- 2) **All programs should compile without reporting any warnings.**
- 3) **Programs without comments will receive a grade of zero.**
- 4) **The assignment must be submitted using the available links in Aula Global. Submitting the assignments by mail is not allowed without prior authorization.**
- 5) **The programs implemented must work in the computers of the informatics lab in the university or the Guernika ([guernika.lab.inf.uc3m.es](http://guernika.lab.inf.uc3m.es)) platform. It is the student responsibility to be sure that the delivered code works correctly in those places.**
- 6) **It is mandatory to follow the input and output formats indicated in each program implemented. In case this is not fulfilled there will be a penalization to the mark obtained.**
- 7) **It is mandatory to implement error handling methods in each of the programs.**
- 8) **Students are expected to submit original work. In case plagiarism is detected between two assignments both groups will fail the continuous evaluation. Additional administrative charges of academic misconduct may be filled.**

**Failing to follow these rules will be translated into zero marks in the affected programs.**

## Appendix – System calls

A system call allows user programs to request services from the operating system. In this sense, system calls can be seen as the interface between the user and kernel spaces. In order to invoke a system call it is necessary to employ the functions offered by the underlying operating system. This section overviews a subset of system calls offered by Linux operating systems that can be invoked in a C program. As any other function, the typical syntax of a system calls follows:

```
status = function (arg1, arg2,.....);
```

### *I/O system calls*

```
int open(const char * path, into flag, ...)
```

The file name specified by *path* is opened for reading and/or writing, as specified by the argument *oflag*; the file descriptor is returned to the calling process.

More information: `man 2 open`

```
int close(int fildes)
```

The `close()` call deletes a descriptor from the per-process object reference table.

More information: `man 2 close`

```
ssize_t read(int fildes, void * buf, size_t nbyte)
```

`read()` attempts to read *nbyte* bytes of data from the object referenced by the descriptor *fildes* into the buffer pointed to by *buf*.

More information: `man 2 read`

```
ssize_t write(int fildes, const void * buf, size_t nbyte)
```

`write()` attempts to write *nbyte* of data to the object referenced by the descriptor *fildes* from the buffer pointed to by *buf*.

More information: `man 2 write`

```
off_t lseek(int fildes, off_t offset, int whence)
```

The `lseek()` function repositions the offset of the file descriptor *fildes* to the argument *offset*, according to the directive *whence*. The argument *fildes* must be an open file descriptor. `Lseek()` repositions the file pointer *fildes* as follows:

- If *whence* is `SEEK_SET`, the offset is set to *offset* bytes.
- If *whence* is `SEEK_CUR`, the offset is set to its current location plus *offset* bytes.

- If whence is SEEK\_END, the offset is set to the size of the file plus *offset* bytes.

More information: `man 2 lseek`

### ***File related system calls***

```
DIR * opendir(const char * dirname)
```

The `opendir()` function opens the directory named by *dirname*, associates a directory stream with it, and returns a pointer to be used to identify the directory stream in subsequent operations.

More information: `man opendir`

```
struct dirent * readdir(DIR * dirp)
```

The `readdir()` function returns a pointer to the next directory entry. It returns NULL upon reaching the end of the directory or detecting an invalid `seekdir()` operation. The *dirent* structure contains a field *d\_name* (`char * d_name`) with the filename and a *d\_type* field (`unsigned char d_type`) with the type of file.

More information: `man readdir`

```
int closedir(DIR * dirp)
```

The `closedir()` function closes the named directory stream and frees the structure associated with the *dirp* pointer, returning 0 on success.



More information: `man closedir`

### ***Manual (man command).***

**man** is a command that formats and displays the online manual pages of the different commands, libraries and functions of the operating system. If a *section* is specified, *man* only shows information about name in that section. Syntax:

```
$ man [section] name
```

A *man* page includes the synopsis, the description, the return values, example usage, bug information, etc. about a *name*. The utilization of *man* is recommended for the realization of all lab assignments. **To exit a man page, press q.**

	<p>Informatics Department BACHELOR'S DEGREE IN COMPUTER SCIENCE AND ENGINEERING Operating Systems(2018-2019) Lab 1 – System calls</p>	
---	---	---

## Bibliography

- El lenguaje de programación C: diseño e implementación de programas Félix García, Jesús Carretero, Javier Fernández y Alejandro Calderón. Prentice-Hall, 2002.
- The UNIX System S.R. Bourne Addison-Wesley, 1983.
- Advanced UNIX Programming M.J. Rochkind Prentice-Hall, 1985.
- Sistemas Operativos: Una visión aplicada Jesús Carretero, Félix García, Pedro de Miguel y Fernando Pérez. McGraw-Hill, 2001.
- Programming Utilities and Libraries SUN Microsystems, 1990.
- Unix man pages (man function)